

# Tutorial for Task 2

Advanced Machine Learning, Fall 2020

## Task 2: Overview

Disease Classification from Image Features:

- Pre-extracted features from image dataset
- Each image is given as a continuous feature vector of size 1000
- 4800 training samples, 4100 test samples
- 3 classes labeled as {0, 1, 2}
- Same class distribution for training and test data

Your Task: Multiclass classification

- Learn a mapping from the features to the set {0, 1, 2}

# Multiclass Classification

We are going to see several ways of doing multiclass classification:


- One-vs-All Strategy
- One-vs-One Strategy
- Multinomial Logistic Regression




# Multiclass Classification: One-vs-All


- For  $C$  classes, we train  $C$  different classifiers
- Classifier  $c$  is constructed using **1** as a label for **class  $c$**  and **0** for the remaining classes.
- After training, the classifier with highest value/confidence is chosen at test time
- **Potential issue:** We might have class imbalance depending on  $C$ .

# Multiclass Classification: One-vs-All

One-vs-all (one-vs-rest):




Class 1:   
Class 2:   
Class 3: 



# Multiclass Classification: One-vs-One

- Create  $\binom{C}{2}$  binary classifiers. Construct classifier  $c_{ij}$  using training data **only** from class  $i$  and  $j$ .
- For each new test point, evaluate all  $\binom{C}{2}$  classifiers
- Usually, the outcome is decided using a voting scheme
- If you get equal votes for some classes, a tie-breaking rule is needed
- Each classifier uses much less data compared to One-vs-All
- **Potential issue:** We might need to train a lot of classifiers, depending on  $C$ .

# Multiclass Classification: One-vs-One



**Fig. 1.** An example of one-versus-one (OVO) decomposition of a three-class problem into three two-class problems.

# Multinomial Logistic Regression

- Reminder: Binary classification loss

$$L(\theta) = - \sum_{i=1}^n y_i \log(p_{\theta}(x_i)) + (1 - y_i) \log(1 - p_{\theta}(x_i))$$

- Extend this to C classes

$$L(\theta) = - \sum_{i=1}^n \sum_{c=1}^C 1_{\{y_i=c\}} \log \left( \frac{\exp(\theta_c^T x_i)}{\sum_{k=1}^C \exp(\theta_k^T x_i)} \right)$$



# Class Imbalance

- Class imbalance refers to the case when the classes are unequally distributed
- Some classes are under or over represented
- In the task 2, you have 4800 training samples
- Only 600 examples for class 0 and 2, 3600 examples for class 1
- Classic accuracy does not work well in this case
- Use Balanced Multiclass Accuracy!

# Balanced Multiclass Accuracy

- Suppose you have the following confusion matrix
- 10 samples from class 1, 70 samples from class 2, 10 samples from class 3

$$M = \begin{bmatrix} 6 & 4 & 0 \\ 4 & 56 & 10 \\ 0 & 3 & 7 \end{bmatrix}$$

- To calculate Balanced Multiclass Accuracy (BMAC), we calculate the true positive rate of each class and take the mean

$$\text{BMAC} = \frac{1}{C} \sum_{c=1}^C \text{TPR}_c$$

$$\text{BMAC} = \frac{\frac{6}{10} + \frac{56}{70} + \frac{7}{10}}{3} = 0.7$$

# How to deal with Class Imbalance?

Possible solutions to deal with class imbalance:

- Oversampling (Adding some instances from the underrepresented class)
- Undersampling (Remove some instances from the overrepresented class)
- Change your loss function using a cost sensitive classifier:

$$L(\theta) = \sum_{i=1}^n l_i(\theta) \quad \longrightarrow \quad L_w(\theta) = \sum_{i=1}^n w_i l_i(\theta)$$

Original cost function puts **equal weight** to every data point

You might want to put **more emphasis** on the **underrepresented** class

For example, in sklearn library, you have the parameter “class\_weight”

# Final Remarks

- You are allowed to use any library
- Pay attention to the class imbalance
- Overfitting to the public score might be dangerous
- Good luck and enjoy the project!